

Accelerate your code with Cython and Numba



- Outline
 - Python C extensions
 - ctypes
 - Cython
 - Numba

what happens when you type python?

- runs a program called "python"
 - Python is actually a C program
 - source code for Python itself is written in C, called CPython
 - alternate implementations in other languages
 - Java: Jython
 - C#: IronPython
 - Python: PyPy
-

why write C extensions?

1. create Python interface to a C library

- example: data acquisition system that only comes with a C interface
- write C extension which allows Python access to that C library, and therefore the acquisition hardware
- writing C extensions by hand is **tedious** and **difficult**
- `DT.c` example
- accessing C libraries is now **much** easier using builtin module `ctypes`
 - lets you directly call any C library function from within Python

2. accelerate your code!

- what's wrong with this code?

```
a = 0
for i in range(1000000000):
    a += 1
```

- code written in C won't have overhead of Python, can *potentially* run much faster

when *not* to accelerate your code:

- when it doesn't matter - don't waste time optimizing code that doesn't take much time anyway


"Premature optimization is the root of all evil" — Donald Knuth



- Pareto principle, i.e. 80/20 rule, i.e. power law distribution
 - probably some small part of your code takes up most of the execution time
- find the part of your code that's the slowest, work on that
 - profile your code to find the slowest parts:
 - `import profile`, `import cProfile`, or in IPython use `prun`

first and best option: numpy!

- built-in functions like `np.mean()` are written in highly-tuned C code
- vectorize! (see upcoming Super Python talk)
- but some complex loops you simply can't vectorize with numpy
- sometimes you need complicated loop(s) that iterates over every entry in an array
- example: gradient ascent clustering algorithm, thousands or millions of points in ND space
 - i. calculate local ND density gradient for every point
 - ii. move each point slightly up its density gradient
 - iii. for each point, check if any other points in ND space are close, if so merge them
 - iv. repeat 1-3 until all remaining points stop moving



Cython: <http://cython.org>

- a much easier way to write C extensions
 - Cython language is superset of Python
 - Python code + static C type declarations
 - evolved from its predecessor called "Pyrex"
1. write your extension in Cython syntax in `.pyx` file
 2. "cython" it to generate C code
 - `cython` is a command line program that converts `.pyx` to `.c`
 3. compile the C code with a compiler
 4. import the compiled C extension from within Python
 - `import myextension`
- `pyximport` compresses these steps into just two: write `.pyx` file, import it in Python
 - in IPython, can use `%%cython` to denote a block of Cython
 - see `cython_example.py`
 - modify Cython example to do RMS instead of mean

the GIL

- Global Interpreter Lock
- allows Python to track references to objects, and therefore do garbage collection and automatic memory deallocation when an object is no longer needed
- makes Python safer, reduces/eliminates memory leaks
- but, the GIL is incompatible with multithreading
 - multithreading: single process, multiple threads w/ shared memory
 - multiprocessing: separate processes w/ separate memory
 - allows taking advantage of multiple CPU cores
 - multithreading can speed up your code without using lots of memory
- turning off the GIL allows you to use multithreading, exposes you to the risks
 - however, if your threads are simple, i.e. your code is "embarrassingly parallel", then it's relatively safe, with big payoff
- put section of Cython code in a `with nogil` block to release the GIL for that code
- or, use `prange()` instead of `range()`
- multithreading in Cython requires OpenMP

Numba: <http://numba.pydata.org>

- newer than Cython, maybe less stable/mature
- Numba is harder to install, so harder to distribute code that relies on Numba
- but, Numba is **much** simpler than Cython. It's close to magic!
- should work out of the box in Anaconda distribution on any platform
- installing separately is a bit tricky:
 - need to install the LLVM compiler "system"
 - LLVM = low-level virtual machine originally
 - like an alternative to GCC, allows for "just in time compiling" similar to Java, but for any language, including C and Python

JIT


- JIT = "Just in time"
 - ▮ Compilation of a function at execution time, as opposed to ahead-of-time compilation.
- "decorate" your Python function with `@jit`, or pass it through `jit()`, and you'll get a faster version
 - what's a decorator?
 - dynamically modifies a function/method
 - little bit analogous to subclassing a class and redefining some methods
 - wraps a function in other code
 - on every call, a decorator gets the original function and its arguments, and can then do whatever it wants with them
 - to stop decorating, remove the `@decorator`, or the `decorator()` call

```
from numba import jit
numba_function = jit(python_function)
```

- <http://numba.pydata.org/numba-doc/dev/user/examples.html>
- see `numba_example.py`
- modify Numba example to do RMS instead of mean

More speed comparisons

- "Numba vs. Cython: Take 2"
 - <https://jakevdp.github.io/blog/2013/06/15/numba-vs-cython-take-2/>
 - comparison of plain python, numpy, cython, Numba, and other acceleration methods,
 - uses example of calculating a 2D numpy array of pairwise distances
 - demos returning a 2D array instead of just a single value



- newer: "Numba vs Cython: How to Choose"
 - <https://eng.climate.com/2015/04/09/numba-vs-cython-how-to-choose/>
- older: "A beginner's guide to using Python for performance computing"
 - <http://scipy.github.io/old-wiki/pages/PerformancePython>

Conclusion

- long loops written in pure Python are inefficient
- optimize the slowest parts of your code first
- various acceleration methods avoid Python loops and do as much as possible directly in C
 - numpy - good for code that can be easily vectorized
 - Cython - fast, flexible, complicated
 - Numba - fast and simple!

